

A bowl of strawberries is shown, with a color gradient overlay that transitions from grayscale on the left to full color on the right. The text is centered over the bowl.


# IMAGE RE-COLORING USING GENERATIVE ADVERSARIAL NETWORKS

Final Project for the 'Neural Networks and Deep Learning' Course

Filippo Canderle, Alberto Merotto, Cecilia Rossi

# INTRODUCTION AND RELATED WORKS

The performed task was **IMAGE RE-COLORING** from Black and White (Grayscale) Images.



# INTRODUCTION AND RELATED WORKS

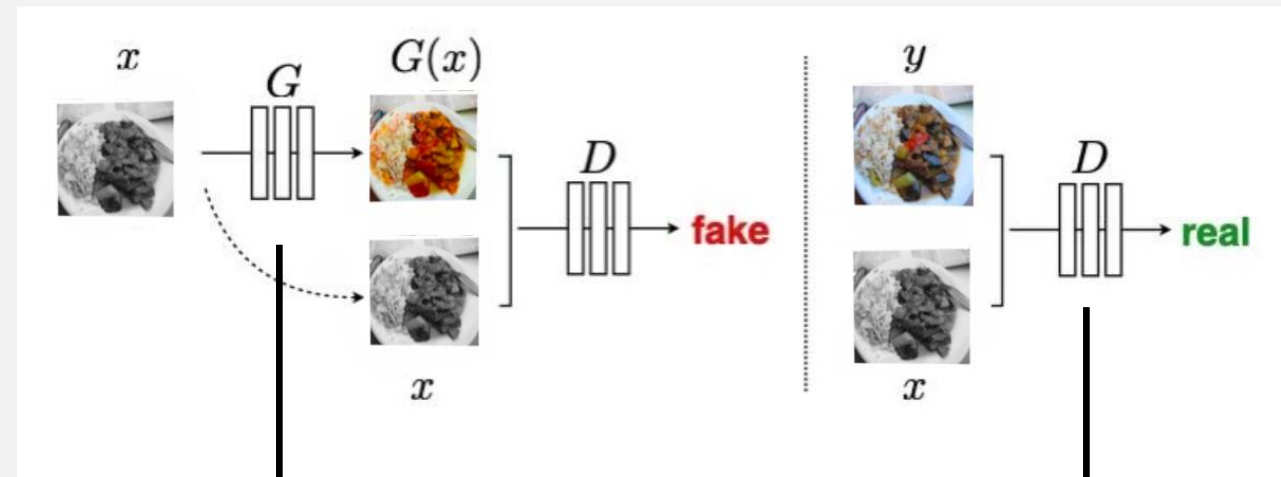
Our basis was the work from Isola, Zhu et al. (2016), 'Image-to-Image Translation with Conditional Adversarial Networks'.

Other related works:

- CycleGAN
- StarGAN
- UNIT

The performed task was **IMAGE RE-COLORING** from Black and White (Grayscale) Images.

We tried to solve this problem relying on a modern approach, based on **(Conditional) Generative Adversarial Network (GAN and cGAN)**:



Generator

**Generates** the **new re-colored images** from the grayscale ones, trying to trick the discriminator mapping from a latent distribution to the target one.

Discriminator

Aims to **classifies the images**, trying to discriminate if they are real (comes from a true distribution) or are generated.

APPROACH

=

in literature is known as:

PIX2PIX



Combination of  
Regression based approaches (MAE) and  
Implicit Density Estimation (BCE)

## 1) Mean Absolute Error (MAE) or L1-Loss

aims to reconstruct pixel wise the image, giving as target the original image.

$$\mathcal{J}(\theta) = \mathbb{E}_{x,y \sim p_{data}} [\|f_{\theta}(x) - y\|]$$

→ not assess joint statistics of the result, and therefore do not measure the very structure that structured losses aim to capture.

## 2) Implicit Density Estimation as Binary Cross Entropy (BCE)

the objective of a Conditional GAN is solving a min-max problem of the following loss:

$$\mathcal{J}(G, D) = \left[ \mathbb{E}_{x \sim p_{data}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [1 - \log D(G(z))] \right]$$

where, G tries to minimize this objective against an adversarial D that tries to maximize it, i.e.

$$G^* = \operatorname{argmin}_G \max_D \mathcal{J}(G, D)$$

## APPROACH

- 1) **Mean Absolute Error (MAE) or L1-Loss**  
aims to reconstruct pixel wise the image, giving as target the original image.

Even though GANs have initially been proposed using BINARY-CROSS ENTROPY (BCE), it has been shown that is not very stable, and that practically they work better with regression losses, as they provide a nicer gradient.

For this reason, we used MAE also for the classification error made by the discriminator:

$$\mathcal{J}(G, D) = \left[ \mathbb{E}_{x, y \sim p_{data}} [\text{L1}(G(x), y)] + \mathbb{E}_{x, y \sim p_{data}} [\log D(x, y)] + \mathbb{E}_{x, y \sim p_{data}} [1 - \log D(1 - D(x, G(x)))] \right]$$

Implicit Density Estimation (BCE)

where, G tries to minimize this objective against an adversarial D that tries to maximize it, i.e.

$$G^* = \underset{G}{\operatorname{argmin}} \max_D \mathcal{J}(G, D)$$



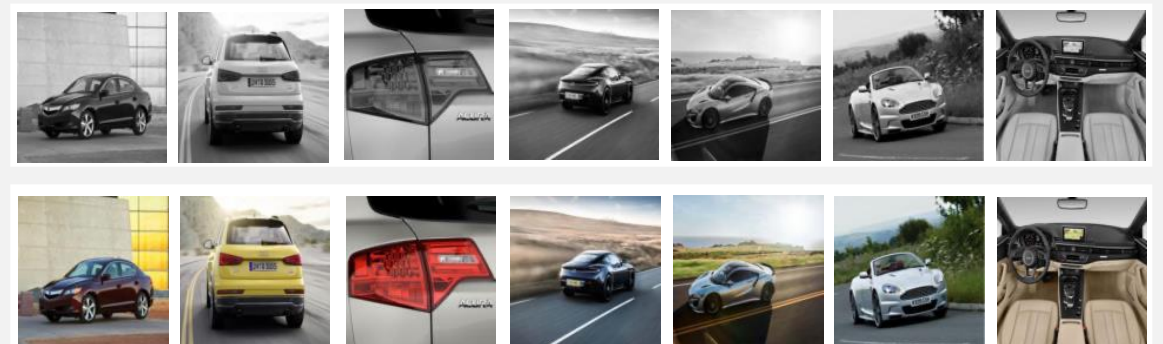
# DATASETS

## COCO DATASET



: a large-scale image recognition dataset for object detection, segmentation, and captioning tasks. We used it as a more Generic Dataset, on which we based our initial analysis.

## CARS DATASET



: a 60k image dataset, obtained from Kaggle. We chose it since we wanted to compare the results obtained on representation of a simple object, that could be designed in slightly different ways and with different color possibilities.

# DATASET PRE-PROCESSING

## RESCALING

Firstly, we resized the images, bringing them to size **128x128 pixels**. In this way:

- We reduced the **computational cost**
- And maintained a **good quality** of the images, preserving visual features

## FLIP

Then, we randomly flipped the images, both **horizontally** and **vertically**, to increase the variability of the data.

## DATALOADERS


Finally, we organized both training and test data into **Python Dataloaders**, created in two different color spaces:

- RGB
- Lab


# LEARNING FRAMEWORK = CONVOLUTIONAL GENERATIVE ADVERSARIAL NETWORK

as said before, it consists of 2 networks:

GENERATOR



DISCRIMINATOR



We tested different popular architectures:


- **U-Net**
- **Autoencoder**
- **ResNet**

We used a fully convolutional neural network, called **Patch Discriminator**



# DISCRIMINATOR

- uses a series of **Convolutional Layer** that, exploiting the fact that convolution has a **local receptive field**, aims to classify the underneath patch of the image, returning a **14x14 patches matrix**.
- The loss is then computed comparing the output of the discriminator with a **matrix of labels**, in which all elements are :
  - 1. associated to real labels;
  - or 0. associated to fake images;



# U-NET

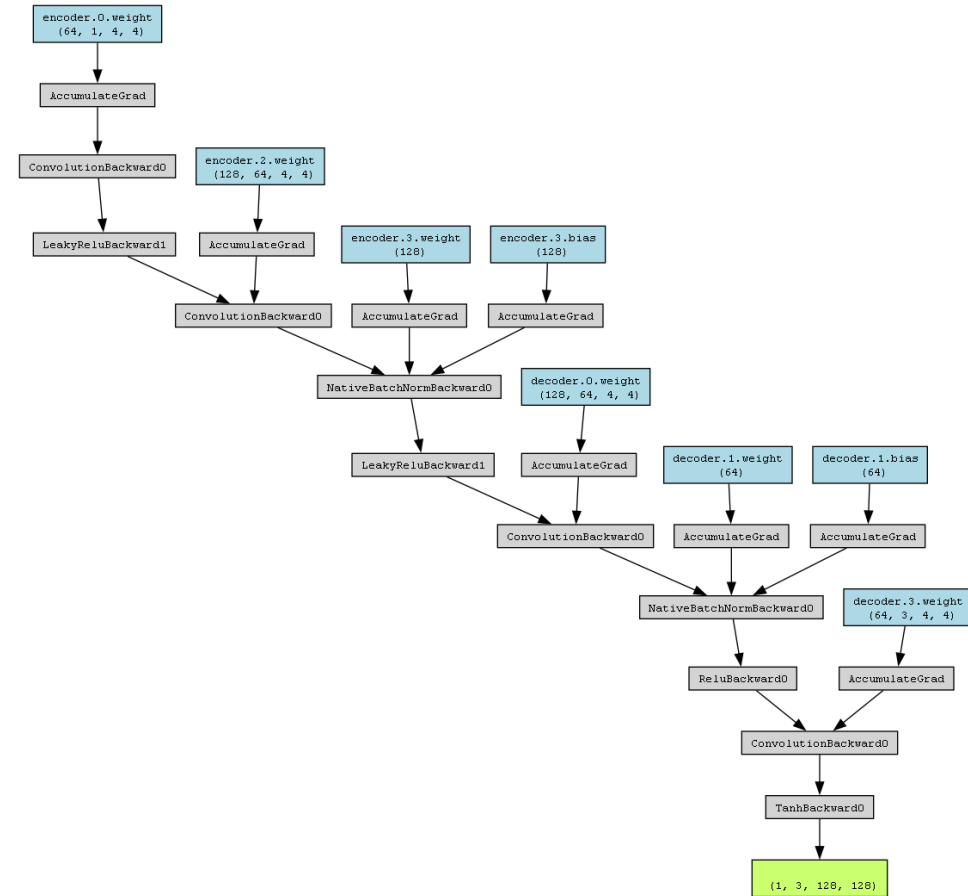


- **Encoder** = 7 down-sampling layers;
- **Decoder** = 7 up-sampling layers;
- Each layers consist of:
  - Linear Convolutional Layer,
  - Batch Normalization,
  - LeakyRelu (encoder) or Relu (decoder) element-wise activation layer.
- **Skip connections**
- **Dropout Layers** in the up-sampling
- Last Convolutional Layer + Tanh activation.

# AUTOENCODER

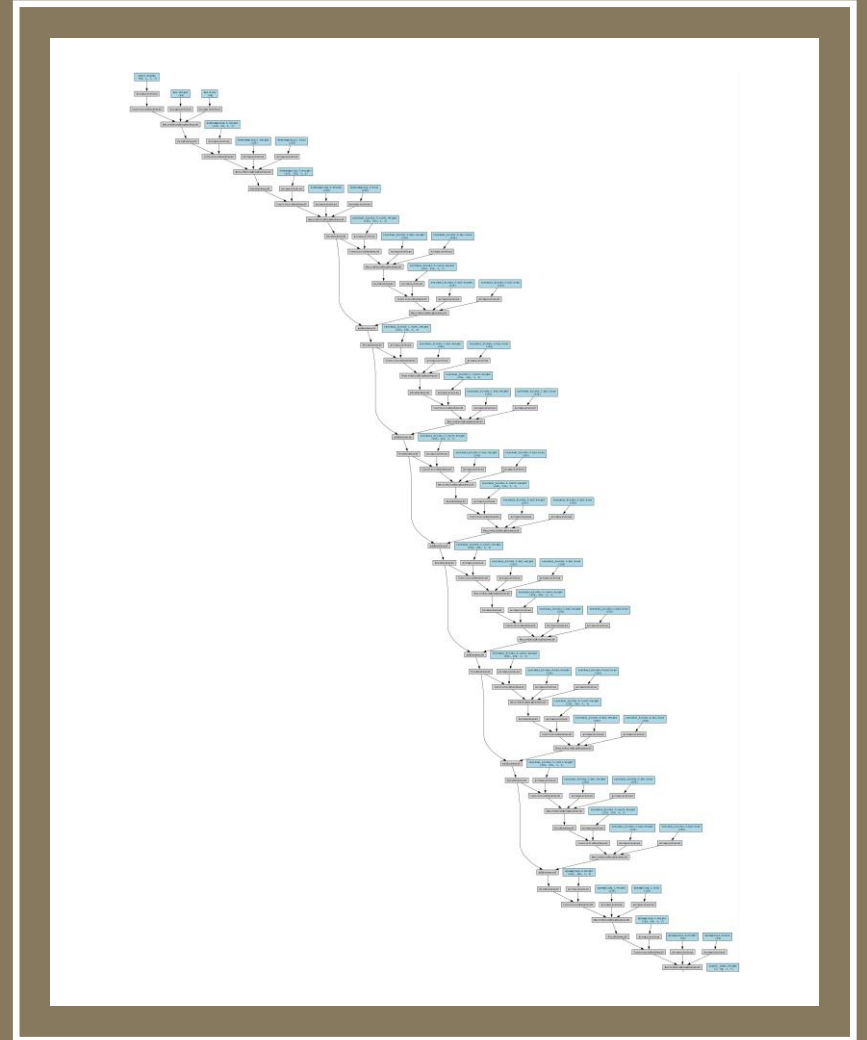
- Very similar to our principal model;
- Main differences:
  - **shallower architecture** (lower numbers of layers both in the down-sampling and up-sampling);
  - use of a **Max Pooling layer**, to effectively decrease spatial dimension;
  - **No Skip Connections;**

We wanted to see if a **less deep network** with the same encoder-decoder structure, but without skip connections, we would have achieved comparable results.



# RESNET

- **Stressing even more the concept of Residual Connection**, in order to reduce the number of parameters.
- It consists of:
  - Initial 7x7 Convolutional Layer with Batch Normalization,
  - Down-sampling module;
  - **Residual Blocks**, to capture and propagate essential image details;
  - Up-sampling module;
  - Final Convolution and Tanh activation.



# CONFIGURATIONS

**U-Net (\*)** on :

- **General** dataset,
- with **10k and 20k** images
- for **15, 30 and 45 epochs**

**GENERAL RESULTS ON U-NET :**  
comparison between training on **different amount of data & different epoch number**

**U-Net** on:

- **Specific** Dataset
- with **10k** images
- for **15, 30 and 45 epochs**

**COMPARISON with SPECIFIC DATASET:**  
(w.r.t. different epoch number)

**U-Net** on:

- **Specific** Dataset
- with **10k LAB** images
- for **45 epochs**

**COMPARISON with LAB COLOR SPACE**

**Autoencoder** on:

- **Specific** Dataset
- with **10k** images
- for **45 epochs**

**COMPARISON between different ARCHITECTURES**

**ResNet** on :

- **Specific** Dataset
- with **10k** images
- for **45 epochs**

(\*) : a cGAN, with the indicated architecture as generator and, as a discriminator, a Patch CNN



## EVALUATION AND RESULTS

- We analysed the plots of the **training losses** of both generator and discriminator.
- In the **evaluation**, we computed the **same losses used during the training on the test set**.
- But, as the majority of “generative” tasks, is very complex to find adequate metrics for the evaluation, in fact:
  - **MAE** (as other regression based metrics) assume a single possible correct output for each pixel, so it's not an effective metric;
  - **Density Based (as BCE)** instead tends to give more importance to the structure of the image, instead of the colorization;

As a consequence of those considerations we relied on the **VISUAL EVALUATION** of the results on a **subset of the test-set**.



# UNET

Trained on  
10k images


B&W input



15 epochs  
training



30 epochs  
training



45 epochs  
training



Original  
image






# UNET

Trained on  
20k images

B&W input



15 epochs  
training




30 epochs  
training



45 epochs  
training



Original  
image





# UNET

Trained on  
10k images  
Car  
Dataset

B&W input



15 epochs  
training



30 epochs  
training



45 epochs  
training



Original  
image



# L A B VS R G B

B&W input



45 epochs  
training Unet  
in RGB color  
space



45 epochs  
training Unet  
in Lab color  
space



Original  
image





# GENERATORS

B&W input



45 epochs  
training  
Autoencoder



45 epochs  
training  
ResNet



45 epochs  
training  
U-Net



Original  
image



## FINAL REMARKS